

Profiling

Profiling measures the performance of a program and can be used to find CPU or memory bottlenecks.

- ▶ `time` A stopwatch
- ▶ `gprof` The GNU (CPU) Profiler
- ▶ `callgrind` Valgrind's CPU profiling tool
- ▶ `massif` Valgrind's memory profiling tool

Timing programs with `time`

- ▶ Just run `time your_program`!
- ▶ Reading `time`'s output:
 - ▶ **Real**: The wall-clock or total time the program took to run.
 - ▶ **User**: The time the program (and libraries) spent executing CPU instructions.
 - ▶ **System**: The time the program spent waiting on system calls (usually I/O).

Profiling with `gprof`

- ▶ You must compile with `g++ -pg program.cpp -o program`.
- ▶ Then, run your program like normal. It will create a file named `gmon.out`.
- ▶ Finally, `gprof program gmon.out` will display profiling statistics!

Understanding `gprof` Output

- ▶ Flat profile: Overview of function usage.
- ▶ Time measures are based on sampling 100 times/second.
- ▶ Function call counts are exact.

Understanding `gprof` Output

- ▶ Flat profile: Overview of function usage.
- ▶ Time measures are based on sampling 100 times/second.
- ▶ Function call counts are exact.
- ▶ Call graph: A listing of which functions called each other.
- ▶ The line with the index entry is the function under consideration.
- ▶ Lines above that are functions that called this function.
- ▶ Lines below that are functions that this function called.

Profiling with `callgrind`

- ▶ As with Memcheck, compile with
`g++ -g program.cpp -o program`
- ▶ Run `valgrind -tool=callgrind ./program`. It will create a file named `callgrind.out.NNNN`.
- ▶ `callgrind_annotate -auto=yes callgrind.out.NNNN` will print some statistics on your program.
- ▶ You can also view the output file directly, although the results are not easy to read.

Understanding `callgrind` Output

- ▶ Callgrind counts instructions executed, not time spent.
- ▶ The annotated source shows the number of instruction executions a specific line caused.
- ▶ Function calls are annotated on the right with the number of times they are called.

Recursion and `callgrind`

- ▶ Recursion can confuse both `gprof` and `callgrind`.
- ▶ The `-separate-recs=N` option to Valgrind separates function calls up to N deep.
- ▶ The `-separate-callers=N` option to Valgrind separates functions depending on which function called them.
- ▶ In general, when you have recursion, the call graph and call counts may be wrong, but the instruction count will be correct.

Profiling with `massif`

- ▶ Compile with `g++ -g program.cpp -o program`
- ▶ Run `valgrind -tool=massif -time-unit=B ./program`.
It will create a file named `massif.out.NNNN`.
- ▶ To get information on stack memory usage as well, include `-stacks=yes` after `-time-unit=B`.
- ▶ `ms_print massif.out.NNNN` will print statistics for you.

Understanding `massif` Output

- ▶ Snapshots: `massif` takes a snapshot of the heap on every allocation and deallocation.
 - ▶ Most snapshots are **plain**. They record only how much heap was allocated.
 - ▶ Every 10th snapshot is **detailed**. These record where memory was allocated in the program.
 - ▶ A detailed snapshot is also taken at peak memory usage.

Understanding `massif` Output

- ▶ Snapshots: `massif` takes a snapshot of the heap on every allocation and deallocation.
 - ▶ Most snapshots are **plain**. They record only how much heap was allocated.
 - ▶ Every 10th snapshot is **detailed**. These record where memory was allocated in the program.
 - ▶ A detailed snapshot is also taken at peak memory usage.
- ▶ The graph: Memory allocated vs. time. Time can be measured in milliseconds, instructions, or bytes allocated.
- ▶ Colons (:) indicate plain snapshots, 'at' signs (@) indicate detailed snapshots, and pounds (#) indicate the peak snapshot.

Understanding `massif` Output

- ▶ Snapshots: `massif` takes a snapshot of the heap on every allocation and deallocation.
 - ▶ Most snapshots are **plain**. They record only how much heap was allocated.
 - ▶ Every 10th snapshot is **detailed**. These record where memory was allocated in the program.
 - ▶ A detailed snapshot is also taken at peak memory usage.
- ▶ The graph: Memory allocated vs. time. Time can be measured in milliseconds, instructions, or bytes allocated.
- ▶ Colons (:) indicate plain snapshots, 'at' signs (@) indicate detailed snapshots, and pounds (#) indicate the peak snapshot.
- ▶ The chart shows the snapshot number, time, total memory allocated, currently-allocated memory, and extra allocated memory.
- ▶ The chart also shows the allocation tree from each detailed snapshot.